

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

train_df = pd.read_csv('mnist_train.csv')
test_df = pd.read_csv('mnist_test.csv')
train_df.head()

  label  1x1  1x2  1x3  1x4  1x5  1x6  1x7  1x8  1x9  ...  28x19
28x20  \
0      5    0    0    0    0    0    0    0    0    0  ...    0
0
1      0    0    0    0    0    0    0    0    0    0  ...    0
0
2      4    0    0    0    0    0    0    0    0    0  ...    0
0
3      1    0    0    0    0    0    0    0    0    0  ...    0
0
4      9    0    0    0    0    0    0    0    0    0  ...    0
0

  28x21  28x22  28x23  28x24  28x25  28x26  28x27  28x28
0      0      0      0      0      0      0      0      0
1      0      0      0      0      0      0      0      0
2      0      0      0      0      0      0      0      0
3      0      0      0      0      0      0      0      0
4      0      0      0      0      0      0      0      0

[5 rows x 785 columns]

test_df.head()

  label  1x1  1x2  1x3  1x4  1x5  1x6  1x7  1x8  1x9  ...  28x19
28x20  \
0      7    0    0    0    0    0    0    0    0    0  ...    0
0
1      2    0    0    0    0    0    0    0    0    0  ...    0
0
2      1    0    0    0    0    0    0    0    0    0  ...    0
0
3      0    0    0    0    0    0    0    0    0    0  ...    0
0
4      4    0    0    0    0    0    0    0    0    0  ...    0
0
```

	28x21	28x22	28x23	28x24	28x25	28x26	28x27	28x28
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0

[5 rows x 785 columns]

train_df.shape

(60000, 785)

y = train_df['label']

x = train_df.drop('label',axis=1)

x_for_test_data = test_df[:]

plt.figure(figsize = (7,7))


some_digit =120

some_digit_image = x.iloc[some_digit].to_numpy().reshape(28,28)

plt.imshow(np.reshape(some_digit_image,(28,28)),cmap=plt.cm.gray)

print(y[some_digit])

2



```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size =
0.30,random_state =42)
```

KNN

```
x_train.shape,y_train.shape,x_test.shape,y_test.shape
((42000, 784), (42000,), (18000, 784), (18000,))

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(x_train,y_train)
```

```
x_train = scaler.transform(x_train)
x_train.shape

(42000, 784)
```

k=5

```
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 5)
classifier.fit(x_train,y_train)
```

```
KNeighborsClassifier()
```

```
y_pred = classifier.predict(x_test)
y_pred
```

```
C:\Users\Dell\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:402: UserWarning: X has feature names, but
KNeighborsClassifier was fitted without feature names
  warnings.warn(
```

```
array([7, 3, 8, ..., 5, 0, 0], dtype=int64)
```

```
from sklearn.metrics import accuracy_score,
classification_report,confusion_matrix
print(accuracy_score(y_test,y_pred))
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

```
0.8108888888888889
```

	precision	recall	f1-score	support
0	0.61	0.98	0.75	1805
1	1.00	0.87	0.93	1994
2	0.91	0.86	0.89	1759
3	0.93	0.73	0.82	1846
4	0.98	0.61	0.75	1726
5	0.98	0.57	0.72	1653
6	0.96	0.92	0.94	1787
7	0.96	0.82	0.89	1937
8	0.51	0.94	0.66	1730
9	0.84	0.77	0.80	1763

accuracy				0.81	18000
macro avg	0.87	0.81	0.81		18000
weighted avg	0.87	0.81	0.82		18000

```
[[1774  0  0  0  0  0  4  0  27  0]
 [  4 1744 26  1  2  1  9  4 201  2]]
```

```
[ 157    0 1512   14    1    0    6    3   63    3]
[ 159    1   60 1343    0    9    1   10  256    7]
[ 278    0    3    2 1047    1   27    4  212  152]
[ 239    0    1   52    0  942   26    1  387    5]
[  88    0    0    0    1    3 1643    0   52    0]
[  52    0   38   15    9    0    0 1597  134   92]
[  60    3   12    8    2    9    3    0 1633    0]
[ 115    1    1    2    2    1    0   41  239 1361]]
```

```
# SVM
```

```
# create a svm classifier
```

```
#from sklearn import svm
```

```
#clf = svm.SVC(kernel='linear', C=10,random_state =0) # linear kernel
```

```
# train the model using training sets
```

```
#clf.fit(x_train,y_train)
```

```
#predict the response for test dataset
```

```
#y_pred_svm = clf.predict(x_test)
```

```
#y_pred_svm
```

```
#from sklearn.metrics import accuracy_score,
```

```
classification_report,confusion_matrix
```

```
#print(accuracy_score(y_test,y_pred_svm))
```

```
#print(classification_report(y_test,y_pred_svm))
```

```
#print(confusion_matrix(y_test,y_pred_svm))
```

k=7

```
from sklearn.neighbors import KNeighborsClassifier
```

```
classifier = KNeighborsClassifier(n_neighbors = 7)
```

```
classifier.fit(x_train,y_train)
```

```
KNeighborsClassifier(n_neighbors=7)
```

```
y_pred = classifier .predict(x_test)
```

```
y_pred
```

```
C:\Users\Dell\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:402: UserWarning: X has feature names, but
KNeighborsClassifier was fitted without feature names
```

```
warnings.warn(
```

```
array([7, 3, 8, ..., 5, 0, 0], dtype=int64)
```

```
from sklearn.metrics import
```

```
accuracy_score,classification_report,confusion_matrix
```

```
print(accuracy_score(y_test,y_pred))
```

```
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

```
0.8176111111111111
```

	precision	recall	f1-score	support
0	0.63	0.99	0.77	1805
1	1.00	0.89	0.94	1994
2	0.92	0.87	0.89	1759
3	0.94	0.75	0.84	1846
4	0.98	0.61	0.75	1726
5	0.99	0.56	0.71	1653
6	0.96	0.92	0.94	1787
7	0.97	0.82	0.89	1937
8	0.51	0.95	0.66	1730
9	0.84	0.79	0.81	1763
accuracy			0.82	18000
macro avg	0.87	0.81	0.82	18000
weighted avg	0.88	0.82	0.82	18000

```
[[1779  0  0  0  0  0  4  0  22  0]
 [  4 1779 27  2  2  0  6  4 169  1]
 [ 139  1 1523 11  2  0  5  2  73  3]
 [ 128  1  44 1385  0  5  0  9 268  6]
 [ 257  0  4  1 1055  0 25  5 228 151]
 [ 231  0  2  45  0 925 31  1 411  7]
 [  82  0  0  0  1  2 1644  0  58  0]
 [  63  0 39 16  8  0  0 1593 130 88]
 [  52  3  9  6  3  5  2  0 1650  0]
 [ 101  1  0  3  3  1  0 33 237 1384]]
```

k=23

```
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 23)
classifier.fit(x_train,y_train)
```

```
KNeighborsClassifier(n_neighbors=23)
```

```
y_pred = classifier .predict(x_test)
y_pred
```

```
C:\Users\Dell\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:402: UserWarning: X has feature names, but
KNeighborsClassifier was fitted without feature names
  warnings.warn(
```

```
array([7, 3, 8, ..., 5, 0, 0], dtype=int64)
```

```

from sklearn.metrics import
accuracy_score, classification_report, confusion_matrix
print(accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

```

0.8309444444444445

	precision	recall	f1-score	support
0	0.68	0.99	0.80	1805
1	1.00	0.91	0.95	1994
2	0.93	0.89	0.91	1759
3	0.94	0.80	0.86	1846
4	0.99	0.61	0.75	1726
5	1.00	0.53	0.70	1653
6	0.95	0.94	0.94	1787
7	0.97	0.83	0.90	1937
8	0.52	0.96	0.67	1730
9	0.84	0.81	0.82	1763
accuracy			0.83	18000
macro avg	0.88	0.83	0.83	18000
weighted avg	0.88	0.83	0.84	18000

```

[[1783  0  0  0  0  0  4  0  18  0]
 [  3 1820 22  3  2  0  9  1 133  1]
 [ 104  0 1574 13  1  0  5  1  59  2]
 [  89  1  37 1480  0  1  1  7 225  5]
 [ 174  0  3  1 1048  0 45  4 273 178]
 [ 234  0  2 65  0 884 28  1 435  4]
 [  78  0  0  0  1  2 1676  0  30  0]
 [  53  0 41 10  3  0  0 1611 132  87]
 [  46  2  8  6  1  1  4  0 1661  1]
 [  73  0  0  3  1  0  1 35 230 1420]]

```

SVM

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

#importing dataset
dataset_train=pd.read_csv('mnist_train.csv') # has 10 classes-- 0 to 9
digits and labels
dataset_test=pd.read_csv('mnist_test.csv')

```

```
print("Shape of training dataset:",dataset_train.shape)
print("Shape of testing dataset:",dataset_test.shape)
```

```
Shape of training dataset: (60000, 785)
Shape of testing dataset: (10000, 785)
```

```
print("Description of training dataset:",dataset_train.describe())
print("Description of test dataset:",dataset_test.describe())
```

```
Description of training dataset:                                label      1x1      1x2
```

1x3	1x4	1x5	1x6	\
count	60000.000000	60000.0	60000.0	60000.0
mean	4.453933	0.0	0.0	0.0
std	2.889270	0.0	0.0	0.0
min	0.000000	0.0	0.0	0.0
25%	2.000000	0.0	0.0	0.0
50%	4.000000	0.0	0.0	0.0
75%	7.000000	0.0	0.0	0.0
max	9.000000	0.0	0.0	0.0

	1x7	1x8	1x9	...	28x19	28x20	\
count	60000.0	60000.0	60000.0	...	60000.000000	60000.000000	
mean	0.0	0.0	0.0	...	0.200433	0.088867	
std	0.0	0.0	0.0	...	6.042472	3.956189	
min	0.0	0.0	0.0	...	0.000000	0.000000	
25%	0.0	0.0	0.0	...	0.000000	0.000000	
50%	0.0	0.0	0.0	...	0.000000	0.000000	
75%	0.0	0.0	0.0	...	0.000000	0.000000	
max	0.0	0.0	0.0	...	254.000000	254.000000	

	28x21	28x22	28x23	28x24	28x25
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
mean	0.045633	0.019283	0.015117	0.0020	0.0
std	2.839845	1.686770	1.678283	0.3466	0.0
min	0.000000	0.000000	0.000000	0.0000	0.0
25%	0.000000	0.000000	0.000000	0.0000	0.0

50%	0.000000	0.000000	0.000000	0.0000	0.0
0.0					
75%	0.000000	0.000000	0.000000	0.0000	0.0
0.0					
max	253.000000	253.000000	254.000000	62.0000	0.0
0.0					

	28x27	28x28
count	60000.0	60000.0
mean	0.0	0.0
std	0.0	0.0
min	0.0	0.0
25%	0.0	0.0
50%	0.0	0.0
75%	0.0	0.0
max	0.0	0.0

[8 rows x 785 columns]

Description of test dataset:				label	1x1	1x2
1x3	1x4	1x5	1x6 \			
count	10000.000000	10000.0	10000.0	10000.0	10000.0	10000.0
10000.0						
mean	4.443400	0.0	0.0	0.0	0.0	0.0
0.0						
std	2.895865	0.0	0.0	0.0	0.0	0.0
0.0						
min	0.000000	0.0	0.0	0.0	0.0	0.0
0.0						
25%	2.000000	0.0	0.0	0.0	0.0	0.0
0.0						
50%	4.000000	0.0	0.0	0.0	0.0	0.0
0.0						
75%	7.000000	0.0	0.0	0.0	0.0	0.0
0.0						
max	9.000000	0.0	0.0	0.0	0.0	0.0
0.0						

	1x7	1x8	1x9	...	28x19	28x20 \
count	10000.0	10000.0	10000.0	...	10000.000000	10000.000000
mean	0.0	0.0	0.0	...	0.179300	0.163600
std	0.0	0.0	0.0	...	5.674149	5.736072
min	0.0	0.0	0.0	...	0.000000	0.000000
25%	0.0	0.0	0.0	...	0.000000	0.000000
50%	0.0	0.0	0.0	...	0.000000	0.000000
75%	0.0	0.0	0.0	...	0.000000	0.000000
max	0.0	0.0	0.0	...	253.000000	253.000000

	28x21	28x22	28x23	28x24	28x25	28x26
28x27 \						
count	10000.000000	10000.0000	10000.0	10000.0	10000.0	10000.0

10000.0						
mean	0.052600	0.0006	0.0	0.0	0.0	0.0
std	2.420004	0.0600	0.0	0.0	0.0	0.0
min	0.000000	0.0000	0.0	0.0	0.0	0.0
25%	0.000000	0.0000	0.0	0.0	0.0	0.0
50%	0.000000	0.0000	0.0	0.0	0.0	0.0
75%	0.000000	0.0000	0.0	0.0	0.0	0.0
max	156.000000	6.0000	0.0	0.0	0.0	0.0

	28x28
count	10000.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

[8 rows x 785 columns]

```
#checking for null and np.nan values
print("Null values in training
dataset:",dataset_train.isnull().sum().head(10))
print("Null values in test
dataset:",dataset_test.isnull().sum().head(10))
```

Null values in training dataset: label 0

1x1	0
1x2	0
1x3	0
1x4	0
1x5	0
1x6	0
1x7	0
1x8	0
1x9	0

dtype: int64

Null values in test dataset: label 0

1x1	0
1x2	0
1x3	0
1x4	0

```

1x5      0
1x6      0
1x7      0
1x8      0
1x9      0
dtype: int64

order = list(np.sort(dataset_train['label'].unique()))
print(order)

[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

# average feature values
round(dataset_train.drop('label', axis=1).mean(), 2)
#drop label column of dataset while calculating mean of all other
columns rounded off upto 2nd

1x1      0.0
1x2      0.0
1x3      0.0
1x4      0.0
1x5      0.0
...
28x24     0.0
28x25     0.0
28x26     0.0
28x27     0.0
28x28     0.0
Length: 784, dtype: float64

# Separating the X and Y variable

y_dataset_train = dataset_train['label']
# Dropping the variable 'label' from X variable
X_dataset_train = dataset_train.drop(columns = 'label') #drop labels
column at index 0
# Printing the size of data
print(X_dataset_train.shape)
print(y_dataset_train.shape)

(60000, 784)
(60000,)

# Normalization-- all values lie between 0 and 255 so by dividing all
values by 255, they will lie between 0 and 1
X_dataset_train = X_dataset_train/255.0
dataset_test = dataset_test/255.0
print("X:",X_dataset_train.shape)
print("Test dataset:",dataset_test.shape)

```

```

X: (60000, 784)
Test dataset: (10000, 785)

from sklearn.preprocessing import MinMaxScaler
#feature scaling independent feature variable X
mm_X = MinMaxScaler()
X_dataset_train = mm_X.fit_transform(X_dataset_train)

#splitting dataset into training set and test set
from sklearn.model_selection import train_test_split
# train test split
X_train, X_test, y_train, y_test = train_test_split(X_dataset_train,
y_dataset_train, test_size = 0.2, random_state = 0)

# Applying SVM model to training dataset
from sklearn.svm import SVC
classifier=SVC(kernel='rbf') #hyperparameter tuning c and gamma values
#linear kernel only gives 91.8% accuracy
#polynomial kernel with degree=3 gives 86% accuracy
#sigmoid kernel gives 91.1% accuracy
#rbf kernel gives 97.816% accuracy so i used rbf kernel
model = classifier.fit(X_train,y_train)

#Predicting results of training dataset
y_pred=classifier.predict(X_test)
print("Predicted output values:",y_pred)

Predicted output values: [3 6 6 ... 5 1 6]

y_pred.flatten()
array([3, 6, 6, ..., 5, 1, 6], dtype=int64)

y_pred.shape
(12000,)

# confusion matrix and accuracy

from sklearn import metrics
from sklearn.metrics import confusion_matrix

accuracy=metrics.accuracy_score(y_true=y_test, y_pred=y_pred)*100
print("Accuracy:", accuracy, "\n")

# cm
print(metrics.confusion_matrix(y_true=y_test, y_pred=y_pred))

# CALCULATING ERROR AND ACCURACY OF PREDICTION MODEL
from sklearn.metrics import
mean_absolute_error,mean_squared_error,r2_score
mse = mean_squared_error(y_test, y_pred)

```

```

print('Mean squared error:',mse)
mae = mean_absolute_error(y_test, y_pred)
print('Mean absolute error:',mae)
r_score = r2_score(y_test, y_pred)
print('Accuracy:',r_score*100)

```

Accuracy: 97.81666666666666

```

[[1195  0  0  0  0  2  6  0  2  0]
 [  0 1372  2  1  2  0  1  0  0  1]
 [  0  1 1142  5  4  1  1  6  5  1]
 [  0  0  9 1169  0 14  0  4 12  0]
 [  0  2  2  0 1132  0  3  4  1  9]
 [  5  0  3 12  4 1035  8  0  5  3]
 [  2  2  3  0  3  5 1173  0  2  0]
 [  0  2  8  1  5  0  1 1199  1 11]
 [  1  5  5  2  2  8  4  0 1161  3]
 [  1  0  0  6 17  5  0 10  6 1160]]

```

Mean squared error: 0.3388333333333333

Mean absolute error: 0.0755

Accuracy: 95.98114200470734

```

from sklearn.metrics import
accuracy_score,classification_report,confusion_matrix
print(accuracy_score(y_test,y_pred))
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

```

0.9781666666666666

	precision	recall	f1-score	support
0	0.99	0.99	0.99	1205
1	0.99	0.99	0.99	1379
2	0.97	0.98	0.98	1166
3	0.98	0.97	0.97	1208
4	0.97	0.98	0.98	1153
5	0.97	0.96	0.97	1075
6	0.98	0.99	0.98	1190
7	0.98	0.98	0.98	1228
8	0.97	0.97	0.97	1191
9	0.98	0.96	0.97	1205
accuracy			0.98	12000
macro avg	0.98	0.98	0.98	12000
weighted avg	0.98	0.98	0.98	12000

```

[[1195  0  0  0  0  2  6  0  2  0]
 [  0 1372  2  1  2  0  1  0  0  1]
 [  0  1 1142  5  4  1  1  6  5  1]
 [  0  0  9 1169  0 14  0  4 12  0]

```

```
[ 0  2  2  0 1132  0  3  4  1  9]
[ 5  0  3 12  4 1035  8  0  5  3]
[ 2  2  3  0  3  5 1173  0  2  0]
[ 0  2  8  1  5  0  1 1199  1 11]
[ 1  5  5  2  2  8  4  0 1161  3]
[ 1  0  0  6 17  5  0 10  6 1160]]
```

Random Forest

```
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.ensemble import RandomForestClassifier
#from sklearn.cross_validation import train_test_split
%matplotlib inline
```

```
data = pd.read_csv('mnist_train.csv')
#test_df = pd.read_csv('mnist_test.csv')
data.head()
```

	label	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	...	28x19
28x20	\											
0	5	0	0	0	0	0	0	0	0	0	...	0
0												
1	0	0	0	0	0	0	0	0	0	0	...	0
0												
2	4	0	0	0	0	0	0	0	0	0	...	0
0												
3	1	0	0	0	0	0	0	0	0	0	...	0
0												
4	9	0	0	0	0	0	0	0	0	0	...	0
0												

	28x21	28x22	28x23	28x24	28x25	28x26	28x27	28x28
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0

```
[5 rows x 785 columns]
```

```
a= data.iloc[3,1:].values
```

```
df_x = data.iloc[:,1:]
```

```
df_y = data.iloc[:,0]
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(df_x,df_y,test_size =
0.30,random_state =42)
```

```
x_train.head()
```

	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	1x10	...	28x19
28x20 \												
918	0	0	0	0	0	0	0	0	0	0	...	0
0												
17141	0	0	0	0	0	0	0	0	0	0	...	0
0												
15558	0	0	0	0	0	0	0	0	0	0	...	0
0												
27327	0	0	0	0	0	0	0	0	0	0	...	0
0												
11606	0	0	0	0	0	0	0	0	0	0	...	0
0												

	28x21	28x22	28x23	28x24	28x25	28x26	28x27	28x28
918	0	0	0	0	0	0	0	0
17141	0	0	0	0	0	0	0	0
15558	0	0	0	0	0	0	0	0
27327	0	0	0	0	0	0	0	0
11606	0	0	0	0	0	0	0	0

```
[5 rows x 784 columns]
```

```
y_train.head()
```

```
918      7
17141    4
15558    5
27327    0
11606    1
Name: label, dtype: int64
```

```
rf = RandomForestClassifier(n_estimators=100)
rf.fit(x_train,y_train)
```

```
RandomForestClassifier()
```

```
pred = rf.predict(x_test)
pred
```

```
array([7, 3, 8, ..., 5, 0, 0], dtype=int64)
```

```
s = y_test.values
count=0
```

```

for i in range(len(pred)):
    if pred[i] == s[i]:
        count = count+1

count
17379

len(pred)
18000

17383/18000.0
0.9657222222222223

from sklearn.metrics import
accuracy_score,classification_report,confusion_matrix
print(accuracy_score(y_test,pred))
print(classification_report(y_test,pred))
print(confusion_matrix(y_test,pred))

```

```
0.9655
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1805
1	0.98	0.98	0.98	1994
2	0.95	0.97	0.96	1759
3	0.96	0.94	0.95	1846
4	0.96	0.97	0.97	1726
5	0.96	0.96	0.96	1653
6	0.98	0.98	0.98	1787
7	0.97	0.96	0.97	1937
8	0.96	0.94	0.95	1730
9	0.95	0.95	0.95	1763
accuracy			0.97	18000
macro avg	0.97	0.97	0.97	18000
weighted avg	0.97	0.97	0.97	18000

```

[[1775  0  2  2  3  4  7  0  9  3]
 [  0 1964 11  7  4  0  2  3  1  2]
 [  4  7 1709  6  5  1  2 16  6  3]
 [  2  1 29 1742  1 28  1 18 14 10]
 [  2  1  2  0 1672  1  6  4  3 35]
 [ 10  4  3 17  2 1585 12  1 13  6]
 [  5  1  3  0  5 11 1756  0  6  0]
 [  1  9 26  0  8  2  0 1868  5 18]
 [  1  8 15 17 15 17  8  3 1633 13]
 [  8  4  6 17 21  7  2 14  9 1675]]

```


Logistic Regression

```
from sklearn.linear_model import LogisticRegression
log_classifier = LogisticRegression(random_state=1000)
log_classifier.fit(x_train,y_train)
import warnings
filterwarnings('ignore')
```

```
C:\Users\Dell\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
LogisticRegression(random_state=1000)
```

```
y_pred_log = log_classifier.predict(x_test)
```

```
y_pred_log
```

```
array([7, 3, 8, ..., 5, 0, 0], dtype=int64)
```

```
print(accuracy_score(y_test,y_pred_log))
```

```
print(classification_report(y_test,y_pred_log))
```

```
print(confusion_matrix(y_test,y_pred_log))
```

```
0.9207777777777778
```

	precision	recall	f1-score	support
0	0.95	0.96	0.96	1805
1	0.96	0.97	0.97	1994
2	0.91	0.89	0.90	1759
3	0.90	0.89	0.90	1846
4	0.93	0.94	0.93	1726
5	0.90	0.87	0.88	1653
6	0.95	0.95	0.95	1787
7	0.94	0.92	0.93	1937
8	0.88	0.89	0.88	1730
9	0.89	0.91	0.90	1763
accuracy			0.92	18000
macro avg		0.92	0.92	18000
weighted avg		0.92	0.92	18000

```

[[1734 1 8 6 9 14 12 4 13 4]
 [ 0 1941 12 9 2 5 1 3 16 5]
 [ 5 20 1574 28 17 8 28 24 45 10]
 [ 9 11 47 1640 1 57 3 17 40 21]
 [ 5 4 8 3 1614 4 15 4 10 59]
 [ 21 9 17 61 10 1440 22 5 60 8]
 [ 11 1 22 1 18 20 1700 1 12 1]
 [ 6 8 25 14 14 3 0 1791 5 71]
 [ 15 28 19 39 14 41 16 5 1538 15]
 [ 10 5 7 17 41 11 0 53 17 1602]]

```

Summary

KNN

K=5 with accuracy 81.1%

K=7 with accuracy 81.7%

K=23 with accuracy 83.1%

Logistic Regression with accuracy 92.1%

Random Forest with accuracy 96.6%

SVM with accuracy 97.8%